

Part I. Warming Up

How do we get all the data in the right place to train a recommendation system, and for real-time inference?

So, you've decided to dive into the world of recommendation systems! Are you hoping to suggest just the right thing based on users' quirky preferences across a vast sea of choices? If so, you've set quite the challenge for yourself! On the surface, these systems might seem straightforward: if User A and User B have similar tastes, then maybe what A likes, B will too. But, as with all things that seem simple, there's a depth that's waiting to be explored.

How do we capture the essence of a user's history and feed it into a model? Where do we stash this model so it's ready to serve up suggestions on the fly? And how do we make sure it doesn't suggest something that steps out of bounds or goes against the business rulebook? Collaborative filtering is our starting point, a guiding light. But there's an entire universe beyond it that makes these systems tick, and together, we're going to navigate it.